RISC-V S-mode Physical Memory Protection (SPMP)

Editor - Dong Du, RISC-V SPMP Task Group

Table of Contents

Preamble

> *This document is in the Development state*
>
> Assume everything can change. This draft specification will change before being accepted as standard, so implementations made to this draft specification will likely not conform to the future standard.

Copyright and license information

This specification is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at creativecommons.org/licenses/by/4.0/.

Copyright 2023 by RISC-V International.

Contributors

The proposed SPMP specifications (non-ratified, under discussion) has been contributed to directly or indirectly by:

- Dong Du, Editor <dd_nirvana@sjtu.edu.cn>
- Bicheng Yang
- Nick Kossifidis
- Andy Dellow
- Manuel Offenberg
- Allen Baum
- Bill Huffman
- Xu Lu
- Wenhao Li
- Yubin Xia
- Joe Xie
- Paul Ku
- Jonathan Behrens
- Robin Zheng
- Zeyu Mi
- Fabrice Marinet

Chapter 1. Introduction

This document describes RISC-V S-mode Physical Memory Protection (SPMP) proposal to provide isolation when MMU is unavailable or disabled. RISC-V based processors recently stimulated great interest in the emerging internet of things (IoT) and automotive devices. However, page-based virtual memory (MMU) is usually undesired in order to meet resource and latency constraints. It is hard to isolate the S-mode OSes (e.g., RTOS) and user-mode applications for such devices. To support secure processing and isolate faults of U-mode software, the SPMP is desirable to enable S-mode OS to limit the physical addresses accessible by U-mode software on a hart.

Chapter 2. S-mode Physical Memory Protection (SPMP)

An optional RISC-V S-mode Physical Memory Protection (SPMP) provides per-hart supervisor-mode control registers to allow physical memory access privileges (read, write, execute) to be specified for each physical memory region.

If PMP/ePMP is implemented, accesses succeed only if both PMP/ePMP and SPMP permission checks pass. The implementation can perform SPMP checks in parallel with PMA and PMP. The SPMP exception reports have higher priority than PMP or PMA exceptions (e.g., an SPMP exception will be raised if the access violates both SPMP and PMP, as will the PMP exception).

SPMP checks will be applied to all accesses whose effective privilege mode is S or U, including instruction fetches and data accesses in S and U mode, and data accesses in M-mode when the MPRV bit in mstatus is set and the MPP field in mstatus contains S or U.

SPMP registers can always be modified by M-mode and S-mode software. SPMP can grant permissions to U-mode, which has none by default. SPMP can also revoke permissions from S-mode.

2.1. Requirements

1) S mode should be implemented 2) `sstatus.SUM` should be WARL.

2.2. S-mode Physical Memory Protection CSRs

Like PMP, SPMP entries are described by an 8-bit configuration register and one XLEN-bit address register. Some SPMP settings additionally use the address register associated with the preceding SPMP entry.

The SPMP configuration registers are packed into CSRs the same way as PMP. For RV32, 16 CSRs, spmpcfg0-spmpcfg15, hold the configurations spmp0cfg-spmp63cfg for the 64 SPMP entries. For RV64, even numbered CSRs (i.e., spmpcfg0, spmpcfg2, …, spmpcfg14) hold the configurations for the 64 SPMP entries; odd numbered CSRs (e.g., spmpcfg1) are illegal. Figures 1 and 2 demonstrate the first 16 entries of SPMP. The layout of the rest entries is similar.

> The terms, entry and rule, are similar to ePMP.
>
> The implementation should decode all SPMP CSRs, and it can modify the number of **writable SPMP entries** while the remaining SPMP CSRs are read-only zero.
>
> The lowest-numbered SPMP entries must be implemented first.

*Figure 1. RV32 SPMP configuration CSR layout*

Figure 2. RV64 SPMP configuration CSR layout

The SPMP address registers are CSRs named spmpaddr0-spmpaddr63. Each SPMP address register encodes bits 33-2 of 34-bit physical address for RV32, as shown in Figure 3. For RV64, each SPMP address encodes bits 55-2 of a 56-bit physical address, as shown in Figure 4. Fewer address bits may be implemented for specific reasons, e.g., systems with smaller physical address space. The number of address bits should be the same for all **writable SPMP entries**. Implemented address bits must extend to the LSB format, except as otherwise permitted by granularity rules in "Address Matching" subsection of PMP.

Figure 3. SPMP address register format, RV32

Figure 4. SPMP address register format, RV64

The layout of SPMP configuration registers is the same as PMP configuration registers, as shown in Figure 5, each field is WARL. The rules and encodings for permission are explained in section 2.4, which resembles the encoding of ePMP (except SPMP does not use locked rules).

1. The S bit marks a rule as **S-mode-only** when set and **U-mode-only** when unset.

2. Bit 5 and 6 are reserved for future use.

3. The A field will be described in the following sections (2.3).

4. The R/W/X bits control read, write, and instruction execution permissions.

7	6	5	4	3	2	1	0
S	Reserved		A		X	W	R
1	2		2		1	1	1

Figure 5. SPMP configuration register format

2.3. Address Matching

The A field in an SPMP entry's configuration register encodes the address-matching mode of the associated SPMP address register. It is the same as PMP/ePMP.

Please refer to the "Address Matching" subsection of PMP in the riscv-privileged spec for detailed information.

> Software may determine the SPMP granularity by writing zero to spmp0cfg, then writing all ones to spmpaddr0, then reading back spmpaddr0. If G is the index of the least-significant bit set, the SPMP granularity is 2 G+2 bytes.

2.4. Encoding of Permissions

SPMP has three kinds of rules: **S-mode-only**, **U-mode-only** and **Shared-Region** rules.

1. An **S-mode-only** rule is **enforced** on Supervisor mode and **denied** on User mode.

2. A **U-mode-only** rule is **enforced** on User modes and **denied**/**enforced** on Supervisor mode depending on the value of `sstatus.SUM` bit:

   - If `sstatus.SUM` is set, a U-mode-only rule is enforced without code execution permission on Supervisor mode to ensure supervisor mode execution protection.

   - If `sstatus.SUM` is unset, a U-mode-only rule is denied on Supervisor mode.

3. A **Shared-Region** rule is enforced on both Supervisor and User modes, with restrictions depending on the `spmpcfg.S` and `spmpcfg.X` bits:

   - If `spmpcfg.S` is not set, the region can be used for sharing data between S-mode and U-mode, yet not executable. S-mode has RW permission to that region, and U-mode has read-only permission if `spmpcfg.X` is not set or RW permission if `spmpcfg.X` is set.

   - If `spmpcfg.S` is set, the region can be used for sharing code between S-mode and U-mode, yet not writeable. S-mode and U-mode have execute permission to the region, and S-mode may also have read permission if `spmpcfg.X` is set.

   - The encoding `spmpcfg.SRWX=1111` is used for backward compatibility, where both S-mode and U-mode have RWX permissions.

4. The encoding `spmpcfg.SRWX=1000` is reserved for future standard use.

The encoding and results are shown in the table:

spmpcfg	S=0			S=1		
	S mode	S mode	U mode	S mode	S mode	U mode
RWX	Sum=0	Sum=1	Sum=x	Sum=0	Sum=1	Sum=x
R - -	Deny	EnforceNoX	Enforce	Enforce	Enforce	Deny
R - X	Deny	EnforceNoX	Enforce	Enforce	Enforce	Deny
- - X	Deny	EnforceNoX	Enforce	Enforce	Enforce	Deny
- - -	Deny	EnforceNoX	Enforce	RSVD		
R W -	Deny	EnforceNoX	Enforce	Enforce	Enforce	Deny
R W X	Deny	EnforceNoX	Enforce	Allow		
- W X	shr RW			shr RX		shr X
- W -	shr RW		shr RO	shr X		

Figure 6. SPMP Encoding Table

Allow: Access allowed.

Deny: Access not allowed.

Enforce: The R/W/X permissions are enforced on accesses.

EnforceNoX: The R/W permissions are enforced on accesses, while the X bit is forced to be zero.

SHR: It is shared between S/U modes with X, RX, RW, or ReadOnly privileges.

RSVD: It is reserved for future use.

SUM bit: The SPMP uses the sstatus.SUM (permit Supervisor User Memory access) bit to modify the privilege with which S-mode loads and stores access to physical memory. The semantics of sstatus.SUM in SPMP is consistent with those in Sv.

MXR bit: The semantics of sstatus.MXR (Make eXecutable Readable) in SPMP is consistent with those in Sv.

2.5. Priority and Matching Logic

Like PMP entries, SPMP entries are also statically prioritized. The lowest-numbered SPMP entry that matches any byte of access (indicated by an address and the accessed length) determines whether that access is allowed or fails. The SPMP entry must match all bytes of access, or the access fails, irrespective of the S, R, W, and X bits.

On some implementations, misaligned loads, stores, and instruction fetches may also be decomposed into multiple accesses, some of which may succeed before an exception occurs. In particular, a portion of a misaligned store that passes the SPMP check may become visible, even if another portion fails the SPMP check. The same behavior may manifest for stores wider than XLEN bits (e.g., the FSD instruction in RV32D), even when the store address is naturally aligned.

1. If the effective privilege mode of the access is M, the access is `allowed`;

2. If the effective privilege mode of the access is S and no SPMP entry matches, the access is `allowed`;

3. If the effective privilege mode of the access is U and no SPMP entry matches, but at least one SPMP entry is implemented, the access is `denied`;

4. Otherwise, the access is checked according to the permission bits in the matching SPMP entry. It is allowed if it satisfies the permission checking with the SRWX encoding corresponding to the access type.

2.6. SPMP and Paging

The table below shows which mechanism to use. (Assume both paged virtual memory and SPMP are implemented.)

satp	Isolation mechanism
satp.mode == Bare	SPMP only
satp.mode != Bare	Paged Virtual Memory only

We do not allow both SPMP and paged virtual memory permissions to be actived at the same time now because: (1) It will introduce one more layer to check permission for each memory access. This issue will be more serious for a guest OS that may have host SPMP and guest SPMP. (2) Paged virtual memory can provide sufficient protection.

That means SPMP is enabled when `satp.mode==Bare` and SPMP is implemented.

> Please refer to Table "Encoding of satp MODE field" in the riscv-privileged spec for detailed information on the satp.MODE field.
>
> If page-based virtual memory is implemented, an SFENCE.VMA instruction with rs1=x0 and rs2=x0 is needed after the SPMP CSRs are written. If page-based virtual memory is not implemented, memory accesses check the SPMP settings synchronously, so no fence is needed. Please refer to hypervisor extension for additional synchronization requirements when hypervisor is implemented.

2.7. Exceptions

Failed accesses generate an exception. SPMP follows the strategy that uses different exception codes for different cases, i.e., load, store/AMO, instruction faults for memory load, memory store/AMO and instruction fetch, respectively.

The SPMP reuses exception codes of page fault for SPMP fault. Because page fault is typically delegated to S-mode, so does SPMP fault, we can benefit from reusing page fault. S-mode software(i.e., OS) can distinguish page fault from SPMP fault by checking satp.mode (as mentioned in 2.6, SPMP and paged virtual memory will not be activated simultaneously). **SPMP proposes to rename page fault to SPMP/page fault for clarity**.

Note that a single instruction may generate multiple accesses, which may not be mutually atomic.

Table of renamed exception codes:

Interrupt	Exception Code	Description
0	12	Instruction SPMP/page fault
0	13	Load SPMP/page fault
0	15	Store/AMO SPMP/page fault

> **ℹ** Please refer to Table "Supervisor cause register (scause) values after trap" in the riscv-privileged spec for detailed information on exception codes.

Delegation: Unlike PMP, which uses access faults for violations, SPMP uses SPMP/page faults for violations. The benefit of using SPMP/page faults is that we can delegate the violations caused by SPMP to S-mode, while the access violations caused by PMP can still be handled by machine mode.

2.8. Context Switching Optimization

With SPMP, each context switch requires the OS to store 64 address registers and 8 configuration registers (RV64), which is costly and unnecessary. So the SPMP proposes an optimization to minimize the overhead caused by context switching.

We add two CSRs called *spmpswitch0* and *spmpswitch1*, which are XLEN-bit read/write registers, as shown in Figure 7. For RV64, only *spmpswitch0* is used. Each bit of this register holds the on/off status of the corresponding SPMP entry. During the context switch, the OS can store and restore spmpswitch as part of the context. An SPMP entry is activated only when both corresponding bits in spmpswitch and A field of spmp[i]cfg are set. (i.e., spmpswitch[i] & spmp[i]cfg.A!=0)

63   0

WARL

Figure 7. SPMP domain switch register format (RV64)

> **ℹ** If the `spmpswitch` is implemented, and `spmpcfg[i].A == TOR`, the entry matches any address y such that spmpaddr[i-1] ≤ y < spmpaddr[i] (irrespective of values of spmpcfg[i-1] and spmpswitch[i-1]).
>
> **The reset state**: On system reset, the `spmpswitch` registers should be zero.

2.9. Access Methods of SPMP CSRs

How SPMP CSRs are accessed depends on whether the `Sscsrind` extension is implemented or not.

Indirect CSR access: The SPMP supports indirect CSR access if the `Sscsrind` extension is implemented. The `Sscsrind` defines 1 select CSR (`siselect`) and 6 alias CSRs (`sireg[i]`). Each combination of `siselect` and `sireg[i]` represents an access to the corresponding SPMP CSR.

siselect number	indirect CSR access of sireg[i]
siselect#1	sireg[1-6] → spmpcfg[0-5]

siselect number	indirect CSR access of sireg[i]
siselect#2	sireg[1-6] → spmpcfg[6-11]
siselect#3	sireg[1-4] → spmpcfg[12-15]
siselect#4	sireg[1-6] → spmpaddr[0-5]
siselect#5	sireg[1-6] → spmpaddr[6-11]
siselect#6	sireg[1-6] → spmpaddr[12-17]
siselect#7	sireg[1-6] → spmpaddr[18-23]
siselect#8	sireg[1-6] → spmpaddr[24-29]
siselect#9	sireg[1-6] → spmpaddr[30-35]
siselect#10	sireg[1-6] → spmpaddr[36-41]
siselect#11	sireg[1-6] → spmpaddr[42-47]
siselect#12	sireg[1-6] → spmpaddr[48-53]
siselect#13	sireg[1-6] → spmpaddr[54-59]
siselect#14	sireg[1-4] → spmpaddr[60-63]
siselect#15	sireg[1-2] → spmpswitch[0-1]

Direct CSR access: SPMP CSRs can be accessed directly with corresponding CSR numbers if the `Sscsrind` extension is not implemented.

> ℹ️ The specific value of `siselect#1-15` will be allocated after review by the Arch Review Committee.
>
> Please refers to the specification of the `Sscsrind` extension for details of indirect CSR access. github.com/riscv/riscv-indirect-csr-access

Chapter 3. Extension of Resource Sharing

Given that PMP and SPMP have similar layout of address/config registers and the same address matching logic. Reusing registers and comparators between PMP and SPMP may be benefitial (in some cases) to save hardware resources. This section introduces the resource sharing extension that can support dynamic reallocation of hardware resource between PMP/ePMP and SPMP. Notably, this extension is not mandatory and a specific implementation can still statically implement the numbers of regions in PMP/ePMP and SPMP.

Implementations can consider PMP/SPMP entries as a resource pool (called PMP Resource). Specifically, one entry of PMP Resource consists the address CSR, configuration CSR, and any micro-architecture states related to an PMP/SPMP entry. A new M-mode CSR called `mpmpimppart` is introduced to control the sharing of PMP Resource between PMP and SPMP.

The CSR has two fields:

1. `pmpimp`: 7-bit, allowing a value of 0—64 to specify the number of PMP entries

2. `spmpimp`: 7-bit, allowing a value of 0—64 to specify the number of SPMP entries

The above two fields allow each of the PMP and SPMP to be of different entries, giving extra flexibility.

Constraints:

1. The values of `pmpimp` and `spmpimp` in `mpmpimppart`, under RV32, can only be a multiple of 4; under RV64, it can only be a multiple of 8. This design avoids sharing the same configuration CSR between S-mode and M-mode.

2. The values of `pmpimp` and `spmpimp` in `mpmpimppart` cannot be larger than 64 (the maximum number of supported PMP/SPMP entries).

3. The sum of `pmpimp` and `spmpimp` cannot be larger than the total number of PMP Resources.

4. Not all physical address bits may be implemented, so the `mpmpimppart` is WARL. This can be utilized for feature discovery.

Addressing:

Notably, the mapping between PMP Resources and PMP/SPMP entries is defined by the specific implementation. PMP entries will be supported contiguously, beginning with the lowest CSR number. Unlike PMP, SPMP entries **may not** start at the lowest CSR number and might include non-contiguous valid entries. For instance, given an implementation with a total of 64 PMP Resource entries, if both `pmpimp` and `spmpimp` are set to 32 during runtime, `PMPResource[0]` to `PMPResource[31]` would map to `PMP[0]` to `PMP[31]`. The remaining entries, `PMPResource[32]` to `PMPResource[63]`, could be mapped as either `SPMP[0]` to `SPMP[31]` or `SPMP[32]` to `SPMP[63]`, illustrating the flexibility in resource sharing design and implementation. If the SPMP entry with lowest CSR number is configured with TOR address-matching mode, zero is used for the lower bound.

Another example is that, an implementation with a total of 96 PMP Resource entries, if `pmpimp` is set to 48 and `spmpimp` is set to 48 during runtime, `PMPResource[0]` to `PMPResource[47]` would map to `PMP[0]` to `PMP[47]`. `PMPResource[48]` to `PMPResource[63]` may map to `SPMP[48]` to `SPMP[63]`, and

`PMPResource[64]` to `PMPResource[95]` may map to `SPMP[0]` to `SPMP[31]`. `SPMP[32]` to `SPMP[47]` are not implemented (from software perspective) in this case.

Re-configuration:

M-mode software can re-configure the entries for PMP and SPMP by modifying the `mpmpimppart` CSR. A re-configuration will be ignored if it will change the PMP resource of a locked PMP entry to an SPMP entry.

> Different mapping implementations have their own trade-offs. For example, mapping `PMPResource[32]` to `PMPResource[63]` onto `SPMP[0]` to `SPMP[31]` could necessitate additional hardware resources to translate the SPMP CSR index to the corresponding `PMPResource`. This might look something like the formula `pmp-resource-index = spmp-index + pmpimp`. Another approach is the non-zero-start mapping. As an example, `PMPResource[32]` to `PMPResource[63]` could be mapped to `SPMP[32]` to `SPMP[63]`, while `SPMP[0]` to `SPMP[31]` are not implemented. In this scenario, the SPMP CSR index mirrors that of the `PMPResources` (when it doesn't exceed 64). A potential drawback is that software must be able to handle scenarios where SPMP doesn't begin at the lowest CSR number. Nevertheless, this is generally acceptable since S-mode software commonly inspects the available SPMP entries by enumerating CSRs at boot time and utilizes the SPMP CSR accordingly based on this probe. Additionally, this design choice becomes even more rational when the Sscsrind extension is implemented and in use. PMP Resouce Sharing extension now supports both mapping methods (and others), and the software only needs to check the usable SPMP entries before using them.

Chapter 4. Summary of Hardware Changes

Item	Changes
CSRs for SPMP address	64 new CSRs
CSRs for SPMP configuration	16 new CSRs for RV32 and 8 for RV64
CSR for domain switch	2 new CSRs for RV32 and 1 for RV64
CSR for resource sharing	1 new CSR
Renamed exception code	***Instruction page fault*** renamed to ***Instruction SPMP/page fault*** ***Load page fault*** renamed to ***Load SPMP/page fault*** ***Store/AMO page fault*** renamed to ***Store/AMO SPMP/page fault***

Chapter 5. Interaction with other proposals

This section discusses how SPMP interacts with other proposals.

J-extension pointer masking proposal: When both PM and SPMP are used, SPMP checking should be performed using the actual addresses generated by PM (pointer masking).

Hypervisor extension: There are extensions (vspmp and hgpmp) in development to support SPMP in guest OS and hypervisors.

Smstateen extension: SPMP adds readable and writable states in S-mode, which can be abused as a covert channel if the OS/hypervisor is not aware of SPMP (thus the states won't be context-switched). It is desired that SPMP occupies a bit in mstateen register of Smstateen extension, which can control supervisor access to SPMP states.